

Help Build the most Advanced SQL Database on Hadoop: **Apache HAWQ**

Lei Chang

lchang@pivotal.io

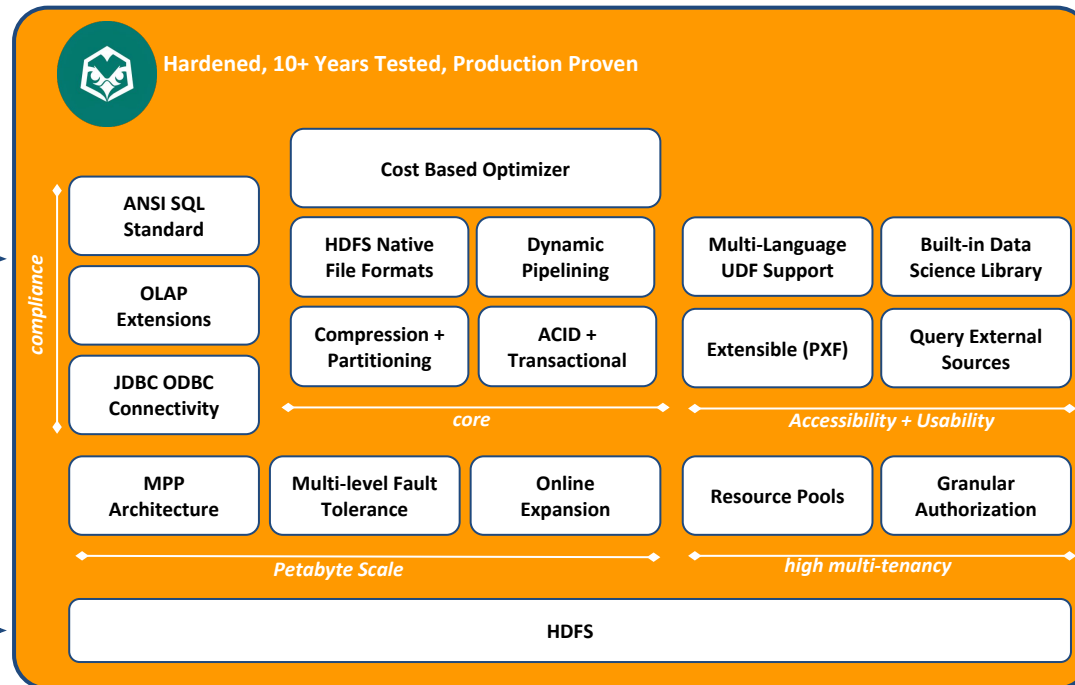
Apachecon 2015, Budapest, Europe





- Introduction
- Architecture
- Query processing
- Catalog service
- Virtual segments
- Interconnect
- Transaction management
- Resource management
- Storage
- How to contribute

HAWQ: A Hadoop Native Parallel SQL Engine



- Leverage Existing Skills & Tools
- Easily Integrate with Other Tools

- Well Integrated with Hadoop Ecosystem



- Discover New Relationships
- Enable Data Science
- Analyze External Sources
- Query All Data Types!



- Manage Multiple Workloads
- Petabyte Scale Analytics
- Sub-second Performance

- Interactive queries
- Scalability
- Consistency
- Extensibility
- Standard compliance
- Productivity




History

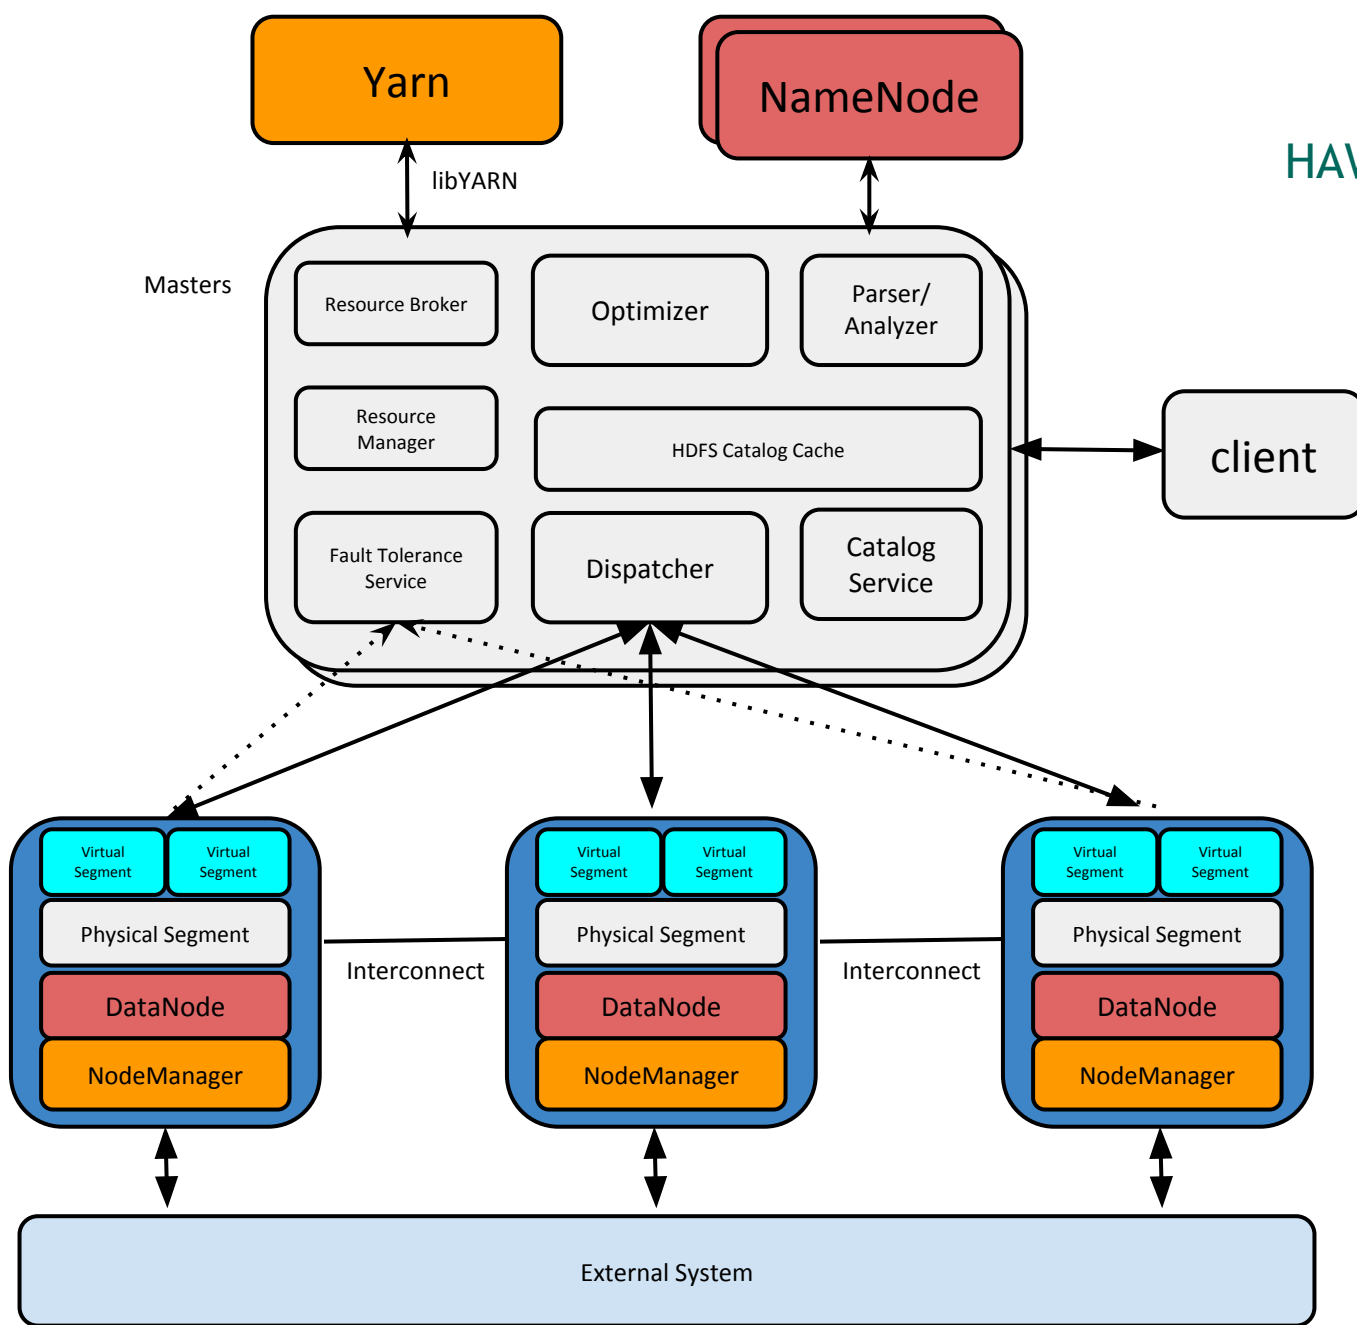


- 2011: Prototype
 - GoH (Greenplum Database on HDFS)
- 2012: HAWQ Alpha
- March 2013: HAWQ 1.0
 - Architecture changes for a Hadoop system
- 2013~2014: HAWQ 1.x
 - HAWQ 1.1, HAWQ 1.2, HAWQ 1.3...
- 2015: HAWQ 2.0 Beta & Apache incubating
 - <http://hawq.incubator.apache.org>

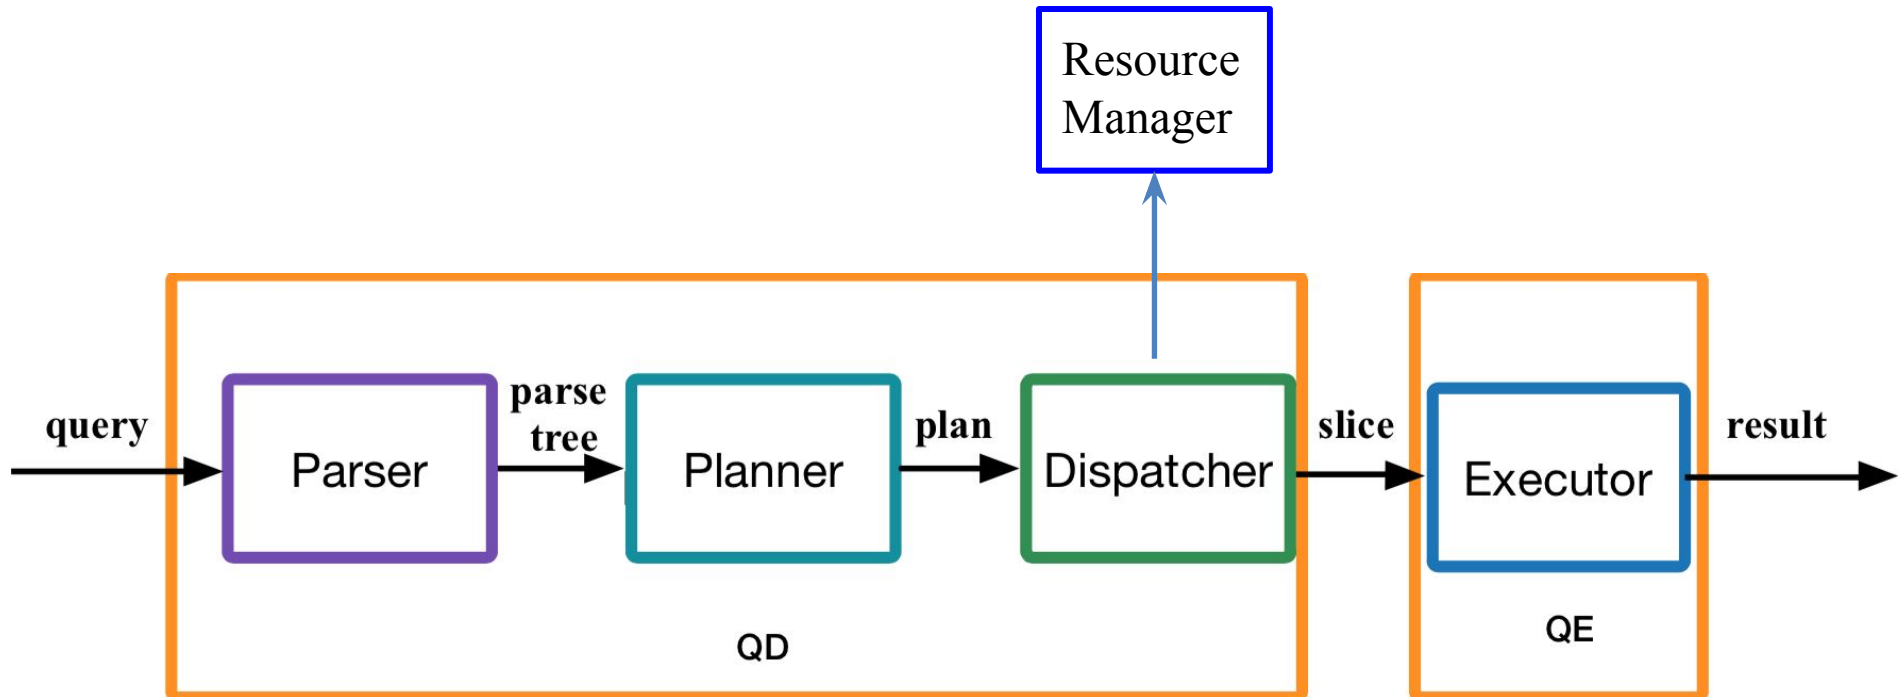
HAWQ components



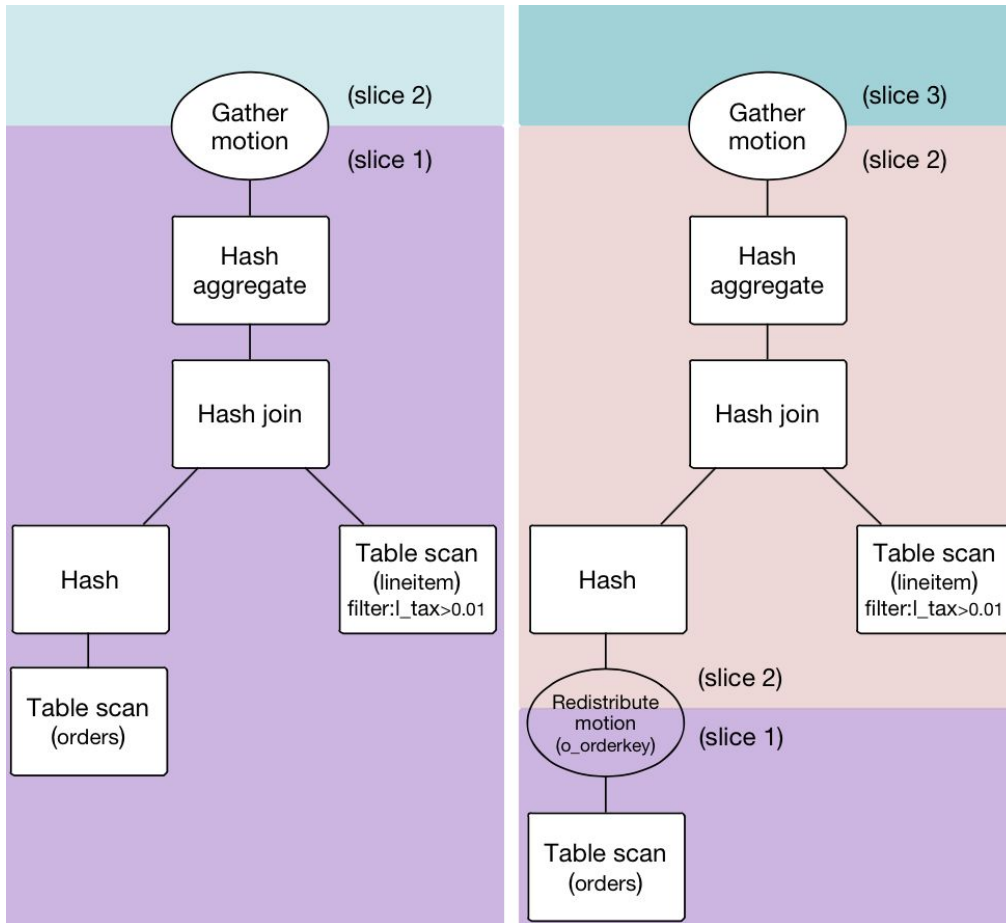
HAWQ Architecture



Basic query execution flow



Plan



Motion:

- Redistribution
- Broadcast
- Gather

```
SELECT l_orderkey, count(l_quantity)
FROM lineitem, orders
WHERE l_orderkey=o_orderkey AND l_tax>0.
GROUP BY l_orderkey;
```


Catalog service

- Currently on master and external in future
- Query language: CaQL
- Basic single-table SELECT, COUNT(), multi-row DELETE, and single row INSERT/UPDATE




Virtual segments & elasticity

- Segment is stateless
- Metadata in Catalog Service
- Self-described plan
- Compression
- Virtual Segment
 - Only one physical segment on each node
 - Ease install and management
 - Ease performance tuning
 - Number of virtual segments
 - Resource available at the query running time
 - The cost of the query
- Node expansion & shrinking




Interconnect




Interconnect


- Two transports: TCP & UDP
- TCP interconnect limitations
 - Limited scale
 - High Latency for connection setup
- Goals of UDP interconnect
 - Reliability
 - Ordering
 - Flow control
 - Performance and scalability
 - Portability



Interconnect state machine



(a) Sender



(b) Receiver

The Design

- Multiple virtual connections share a common socket
- Multi-threaded
- Flow control: flow control window based
- Packet losses and out-of-order packets
- Deadlock elimination



Transaction Management

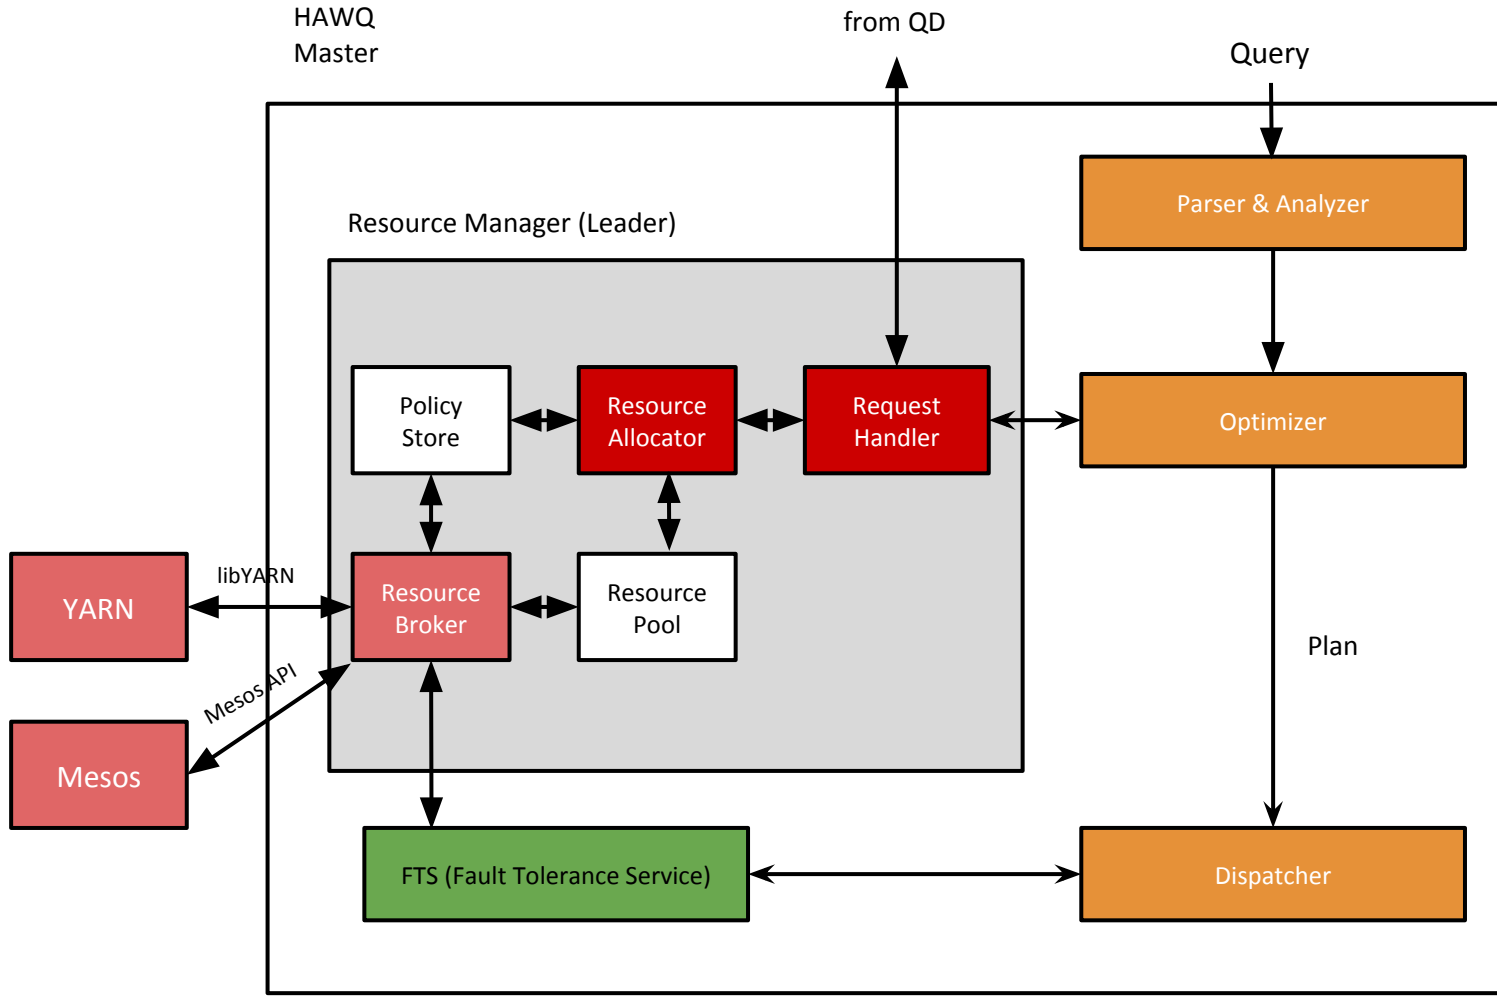
- Snapshot isolation
- Catalog data: MVCC
- User data: Logical EoF & swimming lane
- Truncate to help remove garbage data



Resource manager

- Responsibility
 - Responsible for acquiring from YARN and return resources to YARN
 - Responsible for resource allocation among HAWQ users, queries and operators
- Two level resource management
 - First level: integration with external resource manager (such as YARN...)
 - Second level: HAWQ internal resource manager (user/query level resource management)
 - Third level: operator level
- Hierarchical resource queues
- CPU and memory management (allocation & enforcement)

Resource manager components



Storage

- Row oriented: AO
 - Quicklz, zlib
- PAX like: Parquet
 - Snappy, gzip
- Other format: via PXF



Contributing to HAWQ

- Documentation
 - Wiki
 - Bug reports
 - Bug fixes
 - Features
- Website: <http://hawq.incubator.apache.org/>
 - Wiki: <https://cwiki.apache.org/confluence/display/HAWQ>
 - Repo: <https://github.com/apache/incubator-hawq.git>
 - JIRA: <https://issues.apache.org/jira/browse/HAWQ>
 - Mailing lists: dev/user@hawq.incubator.apache.org


Code contribution process

- Start a JIRA
- Fork a github repo: <https://github.com/apache/incubator-hawq.git>
- Clone your repo to local
- Add the github repo as “upstream”
- Create a feature branch and commit your code
- Start a pull request for code review

Details: <https://cwiki.apache.org/confluence/display/HAWQ/Contributing+to+HAWQ>

Build & Setup

- Build
 - Option 1: Use pre-built docker image: <https://hub.docker.com/r/mayjojo/hawq-dev/>
 - Option 2: Build dependencies by yourself
 - <https://cwiki.apache.org/confluence/pages/viewpage.action?pageId=61320026>
- Setup & Run
 - HDFS (required)
 - YARN (Optional)
 - HAWQ Init/start/stop cluster
 - psql -d postgres



References

- HAWQ website:
 - <http://hawq.incubator.apache.org>
 - <http://pivotal.io/big-data/pivotal-hawq>
- HAWQ research papers
 - Lei Chang et al: [HAWQ: a massively parallel processing SQL engine in hadoop](#). SIGMOD Conference 2014: 1223-1234
 - Mohamed A. Soliman et al: [Orca: a modular query optimizer architecture for big data](#). SIGMOD Conference 2014: 337-348
 - Lyublena Antova et al: [Optimizing queries over partitioned tables in MPP systems](#). SIGMOD Conference 2014: 373-384
 - Amr El-Helw et al: [Optimization of Common Table Expressions in MPP Database Systems](#). PVLDB 8(12): 1704-1715 (2015)

Summary

- HAWQ: A Hadoop Native Parallel SQL Engine
- How to contribute

